

Data manipulations and linear regression

Daniel J. Eck

Background

This lecture is meant to supplement Chapter 2 in your textbook.
We present a brief overview of linear regression.

The `dplyr` package within `tidyverse`

`dplyr` provides comprehensive tools for data manipulations (or wrangling). The main “verbs” are:

- ▶ `select()`: choose from a subset of the columns
- ▶ `filter()`: choose a subset of the rows based on logical criteria
- ▶ `arrange()`: sort the rows based on values of the columns.
- ▶ `mutate()`: add or modify the definitions of the column, and create columns that are functions of existing columns.
- ▶ `summarize()`: collapse a data frame down to a single row (per group) by aggregating vectors into a single value. Often used in conjunction with `group_by()`
- ▶ `left_join()`: add columns from one data set to another, matching observations based on keys.

The pipe operator

The pipe operator `%>%` allows for verbs to be strung in succession so that complicated manipulations can be combined within a single easily digestible sentence.

```
data %>%  
  inner_function() %>%  
  outer_function()
```

Example: Runs differential regression

```
library(tidyverse)
library(Lahman)
data(Teams)
head(Teams, 3)
```

```
##   yearID lgID teamID franchID divID Rank  G  Ghome  W  L DivWin WCWin LgWin
## 1  1871   NA   BS1      BNA  <NA>    3 31    NA 20 10  <NA> <NA>    N
## 2  1871   NA   CH1      CNA  <NA>    2 28    NA 19  9  <NA> <NA>    N
## 3  1871   NA   CL1      CFC  <NA>    8 29    NA 10 19  <NA> <NA>    N
##   WSWin  R  AB  H X2B X3B HR BB SO SB CS HBP SF  RA  ER  ERA CG SHO SV
## 1 <NA> 401 1372 426 70 37 3 60 19 73 16  NA NA 303 109 3.55 22  1  3
## 2 <NA> 302 1196 323 52 21 10 60 22 69 21  NA NA 241 77 2.76 25  0  1
## 3 <NA> 249 1186 328 35 40 7 26 25 18  8  NA NA 341 116 4.11 23  0  0
##   IPouts  HA  HRA  BBA  SOA  E  DP  FP                                name
## 1  828 367  2  42  23 243 24 0.834      Boston Red Stockings
## 2  753 308  6  28  22 229 16 0.829  Chicago White Stockings
## 3  762 346 13  53  34 234 15 0.818  Cleveland Forest Citys
##                                     park attendance BPF PPF teamIDBR teamIDlahman45
## 1                      South End Grounds I          NA 103 98          BOS          BS1
## 2                      Union Base-Ball Grounds      NA 104 102          CHI          CH1
## 3 National Association Grounds                    NA 96 100          CLE          CL1
##   teamIDretro
## 1          BS1
## 2          CH1
## 3          CL1
```



```

dat = Teams %>%
  select(yearID, franchID, W, L, G, AB, H, X2B, X3B, HR, BB, HBP, SF,
         HA, HRA, BBA, SOA, IPouts, FP, R, RA) %>%
  filter(yearID >= 1900) %>%
  replace_na(list(HBP = 0, SF = 0)) %>%
  mutate(RD = (R - RA) / G, X1B = H - (X2B + X3B + HR)) %>%
  mutate(OBP = (H + BB + HBP) / (AB + BB + HBP + SF)) %>%
  mutate(SLG = (X1B + 2*X2B + 3*X3B + 4*HR) / AB) %>%
  mutate(OPS = OBP + SLG) %>%
  mutate(WHIP = 3*(HA + BBA) / IPouts) %>%
  mutate(FIP = 3*(13*BBA + 3*BBA - 2*SOA) / IPouts)
head(dat, 3)

```

```

##   yearID franchID W  L  G   AB    H X2B X3B HR  BB HBP SF    HA HRA BBA SOA
## 1   1900      LAD 82 54 141 4860 1423 199  81 26 421  81  0 1370  30 405 300
## 2   1900      ATL 66 72 142 4952 1403 163  68 48 395  45  0 1263  59 463 340
## 3   1900      CHC 65 75 146 4907 1276 202  51 33 343  65  0 1375  21 324 357
##   IPouts  FP  R  RA      RD  X1B      OBP      SLG      OPS  WHIP
## 1   3677 0.948 816 722  0.6666667 1117 0.3590078 0.3831276 0.7421354 1.448191
## 2   3721 0.953 778 739  0.2746479 1124 0.3418027 0.3727787 0.7145813 1.391561
## 3   3813 0.933 635 751 -0.7945205  990 0.3168391 0.3421643 0.6590034 1.336743
##
##           FIP
## 1 0.8199619
## 2 1.1900027
## 3 0.4177813

```


Note: other packages may contain functions with the same name as those in `dplyr`. For example, the `MASS` package also contains a `select` function.

In the event that you have both `dplyr` and `MASS` loaded in an R session, you can access `dplyr`'s `select` function using `dplyr::select`

```
dat = Teams %>%
  dplyr::select(yearID, franchID, W, L, G, AB, H, X2B, X3B, HR, BB, HBP, SF,
               HA, HRA, BBA, SOA, IPouts, FP, R, RA) %>%
  filter(yearID >= 1900) %>%
  replace_na(list(HBP = 0, SF = 0)) %>%
  mutate(RD = (R - RA) / G, X1B = H - (X2B + X3B + HR)) %>%
  mutate(OBP = (H + BB + HBP) / (AB + BB + HBP + SF)) %>%
  mutate(SLG = (X1B + 2*X2B + 3*X3B + 4*HR) / AB) %>%
  mutate(OPS = OBP + SLG) %>%
  mutate(WHIP = 3*(HA + BBA) / IPouts) %>%
  mutate(FIP = 3*(13*BBA + 3*BBA - 2*SOA) / IPouts)
```

Baseball is a game of offense, pitching, and defense. Let's see how well runs differential per game is explained by:

- ▶ **OPS**: on base percentage plus slugging percentage
- ▶ **WHIP**: walks and hits allowed divided by innings pitched
- ▶ **FP**: fielding percentage

using a linear regression model

```
m = lm(RD ~ OPS + WHIP + FP, data = dat)
```

Regression review

Regression model:

$$y = \beta_1 x_1 + \cdots + \beta_p x_p + \varepsilon; \quad \varepsilon \sim N(0, \sigma^2),$$

where we usually specify a model intercept by setting $x_1 = 1$.

Can also write in vector notation:

$$y = \mathbf{x}'\beta + \varepsilon; \quad \varepsilon \sim N(0, \sigma^2),$$

where $\mathbf{x}, \beta \in \mathbb{R}^p$.

Either way, this model relies on a few assumptions:

- ▶ a linear relationship is present
- ▶ errors are independent and identically distributed
- ▶ errors are normally distributed mean 0 and common variance σ^2

Regression review

Remember that linear regression is about modeling a conditional expectation, the scattering of points is noise. Interest is in

$$E(y|\mathbf{x}) = \mathbf{x}'\beta,$$

where it is important to choose the variables comprising \mathbf{x} and to be able to defend those choices.

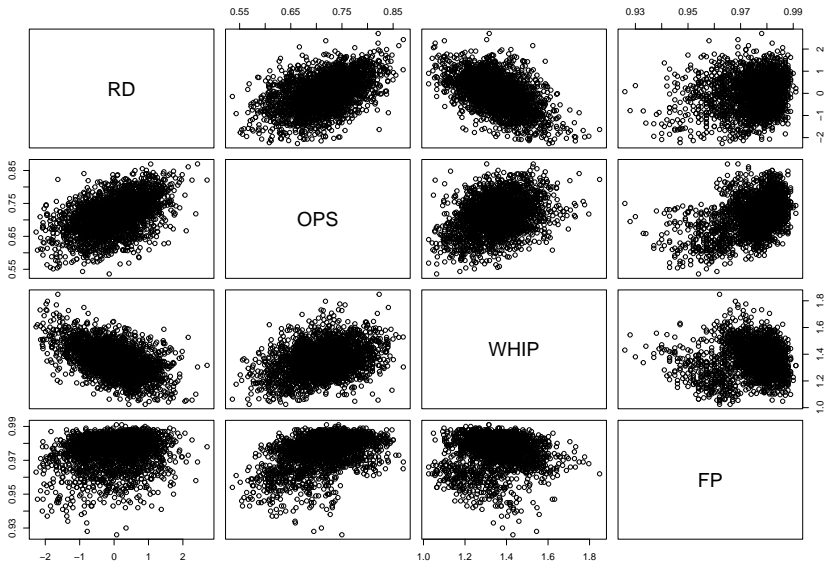
Yes, baseball IS a game of offense, pitching, and defense.

`summary(m)`

```
##
## Call:
## lm(formula = RD ~ OPS + WHIP + FP, data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.4707 -0.2169 -0.0061  0.2103  1.4758
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  13.71282    0.78206   17.53  <2e-16 ***
## OPS          11.79734    0.15538   75.92  <2e-16 ***
## WHIP         -5.40521    0.06403  -84.42  <2e-16 ***
## FP          -15.19191    0.83208  -18.26  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3484 on 2636 degrees of freedom
## Multiple R-squared:  0.7891, Adjusted R-squared:  0.7889
## F-statistic: 3288 on 3 and 2636 DF, p-value: < 2.2e-16
```


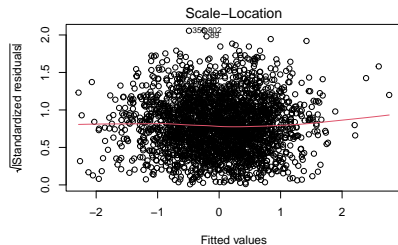

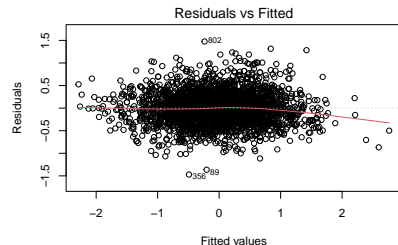
Linearity holds.

```
pairs(dat %>% select(RD, OPS, WHIP, FP))
```



Normality of mean-zero errors with constant variance holds.
Although a slightly heavy right tail is observed in the residuals.

```
par(mfrow = c(2,2))  
plot(m)
```



The compiled fractions should be roughly 68% and 95% if errors are truly normal with common variance.

```
library(broom)
dat_aug = augment(m, data = dat)
dat_aug %>%
  mutate(rmse = sqrt((mean(.resid^2)))) %>%
  summarise(N = n(),
            within_1rmse = sum(abs(.resid) < rmse),
            within_2rmse = sum(abs(.resid) < 2 * rmse)) %>%
  mutate(within_1rmse_pct = within_1rmse / N,
         within_2rmse_pct = within_2rmse / N)

## # A tibble: 1 x 5
##       N within_1rmse within_2rmse within_1rmse_pct within_2rmse_pct
##   <int>      <int>      <int>          <dbl>          <dbl>
## 1  2640      1872      2496          0.709          0.945
```


We will suppose that independence holds, or that any violations of this assumption that may be present in this data do not materially effect our overall conclusions.

A saturated model (one parameter per observation) does not fit the data better than our model with three variables and an intercept.

```
# likelihood ratio test of fitted model vs a saturated model  
m_glm = glm(RD ~ OPS + WHIP + FP, data = dat)  
pchisq(m_glm$deviance, m_glm$df.residual, lower = FALSE)
```

```
## [1] 1
```

Thus we have a well-fitting simple and useful model that provides satisfactory dimension reduction.

Investigate large residuals

```
dat_aug %>% filter(abs(.resid) >= 1) %>%  
  select(yearID, franchID, W, RD, OPS, WHIP, FP, .resid, .fitted) %>%  
  mutate(across(4:9, round, 3)) %>%  
  arrange(desc(.resid))
```

```
## Warning: There was 1 warning in `mutate()`.  
## i In argument: `across(4:9, round, 3)`.  
## Caused by warning:  
## ! The `...` argument of `across()` is deprecated as of dplyr 1.1.0.  
## Supply arguments directly to `.fns` through an anonymous function instead.  
##  
## # Previously  
## across(a:b, mean, na.rm = TRUE)  
##  
## # Now  
## across(a:b, \(x) mean(x, na.rm = TRUE))  
  
## # A tibble: 26 x 9  
##   yearID franchID      W      RD      OPS      WHIP      FP .resid .fitted  
##   <int> <fct>      <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1  1949 NYY         97  1.24  0.758  1.49  0.977  1.48 -0.237  
## 2  1936 NYY        102  2.15  0.861  1.53  0.973  1.32  0.84  
## 3  1939 NYY        106  2.70  0.821  1.32  0.978  1.28  1.42  
## 4  1950 BOS         94  1.45  0.846  1.59  0.981  1.24  0.213  
## 5  1949 BOS         96  1.48  0.8    1.48  0.98    1.21  0.269  
## 6  1948 BOS         96  1.21  0.779  1.48  0.981  1.19  0.02  
## 7  1935 DET         93  1.67  0.798  1.44  0.979  1.19  0.485  
## 8  1943 CIN         87  0.419 0.653  1.34  0.98    1.12 -0.697  
## 9  1950 NYY         98  1.44  0.804  1.48  0.979  1.11  0.328  
## 10 1938 DET         84  0.432 0.768  1.59  0.976  1.07 -0.634  
## # i 16 more rows
```

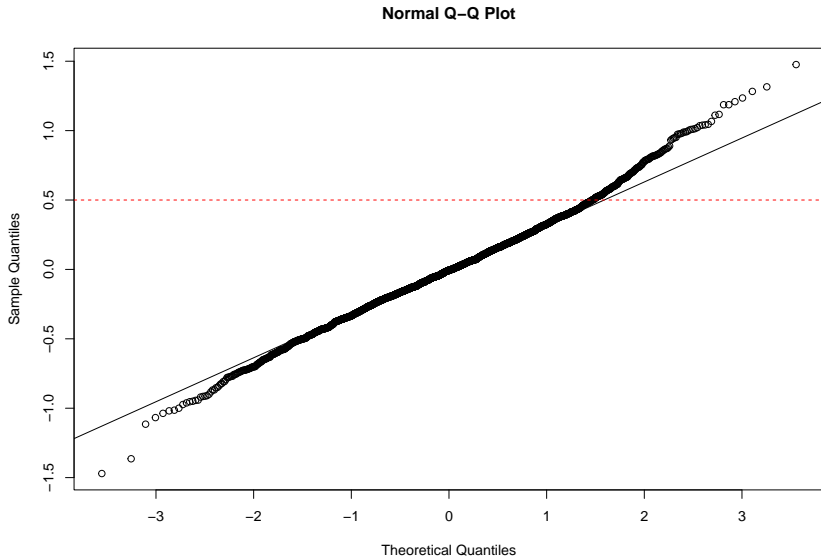
Investigate large fitted values

```
dat_aug %>% filter(.fitted >= 2) %>%  
  select(yearID, franchID, W, RD, OPS, WHIP, FP, .resid, .fitted)
```

```
## # A tibble: 5 x 9  
##   yearID franchID      W      RD      OPS      WHIP      FP .resid .fitted  
##   <int> <fct>      <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1  1927  NYN      110  2.43  0.870  1.30  0.969  0.221  2.20  
## 2  2019  HOU      107  1.73  0.848  1.13  0.988 -0.868  2.60  
## 3  2019  LAD      106  1.69  0.810  1.10  0.982 -0.705  2.39  
## 4  2020  LAD       43  2.27  0.821  1.06  0.982 -0.500  2.77  
## 5  2022  LAD      111  2.06  0.775  1.05  0.986 -0.152  2.21
```

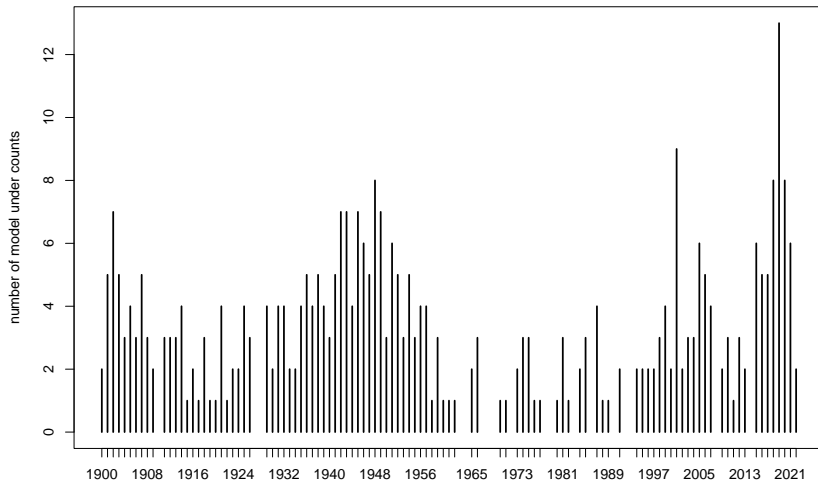
A closer look at problems with fit

```
qqnorm(resid(m)); qqline(resid(m))  
abline(a=0.5, b=0, lty = 2, col = "red")
```



A closer look at problems with fit

```
plot(table(dat_aug %>% filter(abs(.resid) >= 0.5) %>%  
  pull(yearID)), ylab = "number of model under counts")
```



League conditions change over time

